

Thin Locks: Featherweight Synchronization for Java

D. Bacon¹ R. Konuru¹ C. Murthy¹ M. Serrano¹

Presented by: Calvin Hubble²

¹IBM T.J. Watson Research Center

²Department of Computer Science
University of CA, San Diego

16th November 2005

Outline

Introduction

- Problem Formulation

Thin Locks

- Implementation

- Locking Algorithms

Evaluation

- Benchmarks

- Conclusions

Questions

Outline

Introduction

Problem Formulation

Thin Locks

Implementation

Locking Algorithms

Evaluation

Benchmarks

Conclusions

Questions

Introduction

- ▶ Java synchronization is a double-edged word
 - ▶ Java has threads and synchronized methods
 - ▶ Synchronization is “dog slow”
- ▶ Stuck with a tradeoff
 - ▶ Bad Performance, Safe Code
 - ▶ Good performance, bug-prone code
- ▶ Can we modify Java to be faster yet still thread-safe to the everyday programmer?

Problem


- ▶ Because Java is an explicitly multi-threaded language, general-purposes libraries are thread-safe
 - ▶ Non-trivial public methods of standard utility classes like **Vector** or **Hashtable** are synchronized
 - ▶ Example: Library call to set a bit in a bit vector:
 - ▶ 50 instructions to lock and unlock the object
 - ▶ 10 instructions method call overhead
 - ▶ 5 instructions to actually set the bit
 - ▶ Locking overhead frequently 25 – 50%
 - ▶ Even in single-threaded applications!!

Locking Scenarios by Frequency

1. Locking an unlocked object
2. Locking an object already locked by current thread a small number of times (shallow nested locking)
3. Locking an object already locked by the current thread many times (deeply nested locking)
4. Being the first to queue on a locked object
5. Trying to lock an object with a queue

Measurements: median of 80% of all lock operations are on unlocked objects, and nesting is very shallow.

Locking Frequency



Goal

Goal: a locking algorithm with very low overhead for single-threaded programs, but with excellent performance in the presence of multithreading and contention.

Outline

Introduction

Problem Formulation

Thin Locks

Implementation

Locking Algorithms

Evaluation

Benchmarks

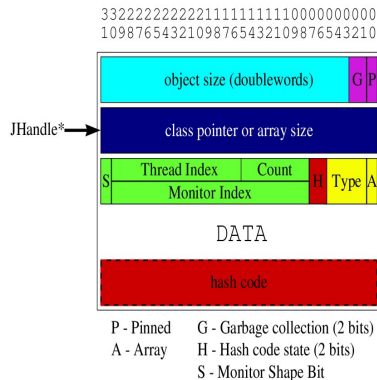
Conclusions

Questions

Thin Locks


- ▶ Assume pre-existing heavy-weight locking system
 - ▶ “Fat Locks”
- ▶ Thin Locks - a lightweight system for 2 most common cases
 1. Object is unlocked
 2. Shallow nested locking
- ▶ Locks are defaulted to thin and inflated if needed
- ▶ Once a lock is inflated, it can never be defaulted

- ▶ Reserve 24 bits in the header of each object for a thin lock
 - ▶ “Obtained 24 free bits using various encoding techniques for the other values typically stored in the header”
- ▶ First bit: Monitor shape lock
 - ▶ 0 - denotes lock is “thin”
 - ▶ 1 - denotes lock is “fat”




When Lock is Thin

- ▶ Lock Structure
 - ▶ Monitor Shape bit - 0
 - ▶ Next 15 bits - Thread Identifier
 - ▶ Last 8 bits - Nested lock count (+1)
- ▶ Maximum of 255 nested locks



When Lock is Fat

- ▶ Lock Structure
 - ▶ Monitor Shape bit - 1
 - ▶ Next 23 bits - index of fat lock



Assumptions

- ▶ Hardware support
 - ▶ Used **compare-and-swap**
 - ▶ `CMP&SWP(addr, old, new)` - If contents of `addr` == `old` value, store new value and return true, otherwise return false
- ▶ key invariant: The lock field is never modified by any thread except the current “owner”.

Locking without Contention

- ▶ Initially - lock field is 0, thread *A* wishes to lock.
- ▶ Algorithm:
 - ▶ compare-and-swap lock word
 - ▶ “Old” value: High 24-bits masked to 0
 - ▶ “New” value: monitor shape 0, thread index *A*, count 0
 - ▶ If succeeds, object was not locked by another thread and we now own lock

Unlocking without Contention (no nesting)

- ▶ Algorithm
 - ▶ Construct “old” value: monitor shape 0, thread index A , count 0
 - ▶ Read lock word and check if compares to old value, if so replace with all 0s
- ▶ Does not need compare-and-swap since no other thread can modify lock if we own it

Nested Locking and Unlocking

- ▶ Locking
 - ▶ Compare-and-swap (from before) will fail
 - ▶ If (monitor shape == 0) and (thread index == A) and (count < 255)
 - ▶ Increment count field - If count overflows then inflate lock
- ▶ Unlocking
 - ▶ Similar to above only decrement lock-count

Locking with Contention

B tries to acquire a lock held by A

- ▶ B 's compare-and-swap will fail
- ▶ B 's check that B owns the lock (nested lock) will fail
- ▶ B needs to force a transition from thin to fat
 - ▶ B enters a spin-locking loop
 - ▶ Once A unlocks, B will obtain
 - ▶ B creates a fat lock, assigns monitor index to new monitor
 - ▶ B changes monitor shape bit to 1

Outline

Introduction

Problem Formulation

Thin Locks

Implementation

Locking Algorithms

Evaluation

Benchmarks


Conclusions

Questions

Setup


- ▶ JDK111
 - ▶ Straightfowrard port of Sun's JDF 1.1.1 to AIX
- ▶ IBM112
 - ▶ IBM's 1.1.2 version of the JDK for AIX
 - ▶ Assumes that most apps have a small number of heavily used locks
 - ▶ Pre-allocates 32 “hot locks”
 - ▶ Suffers when a large number of locks are used
- ▶ ThinLock
 - ▶ Implementation of thin locks in JDK 1.1.2 for IBM's AIX OS

Micro Benchmarks



Benchmarks

Macro Benchmarks



Conclusions

- ▶ Efficient
 - ▶ 5-10 instructions to lock/unlock object
 - ▶ no increase in object size
- ▶ Good speedups
- ▶ Portable
 - ▶ All architectures offer some locking primitive

Outline

Introduction

Problem Formulation

Thin Locks

Implementation

Locking Algorithms

Evaluation

Benchmarks

Conclusions

Questions

Questions

anyone? anyone? Bueller?